

Quantifying Model, Label, and Analyst Contributions to Variance in Real-Time Sepsis Onset Prediction

Abdul Ahad

M.Sc. in Computing

Department of Computer Science & Applied Physics

Atlantic Technological University (ATU), Galway

I. Introduction

- **Sepsis** is a life-threatening condition caused by the body's extreme response to an infection [1].
- It causes **about 11 million deaths every year**, nearly **1 in 5 deaths worldwide** [2].
- Delaying treatment by even a few hours increases the risk of death [3].
- **More than 90 AI models** have already been developed to predict sepsis [4].

Existing models have improved prediction, but important methodological questions remain unanswered.

II. Background and Related Work

Metric	Internal	External
Median AUROC	0.811	0.783
Median Utility Score	+0.381	-0.164

Four key limitations of the conventional approach:

1. **The sepsis label can change.** Different Sepsis-3 criteria create very different patient groups (867–2,178 cases) [5].
2. **The data split can give unrealistic results.** Performance drops significantly with temporal splits [6].
3. **AUROC does not show the full picture.** A model can maintain good AUROC while clinical usefulness decreases [4].
4. **The label may represent doctor decisions, not patient deterioration** [7].

III. Research Question and Hypotheses

Research Question

In real-time sepsis prediction using MIMIC-IV [8], how much of the model's performance difference comes from (a) the choice of AI model, and (b) the researcher's decisions, including how sepsis is defined, how data is split, when prediction is made, and how performance is measured?

#	Hypothesis	Threshold
H1	Label choice affects AUROC $\geq$ model choice	label $\geq$ model spread
H2	Feature importance changes across labels	≥ 1 sign change
H3	Removing post-treatment info lowers AUROC	$\Delta \geq 0.03$
H4	Temporal $<$ random testing	$\Delta \geq 0.02$
H5	Utility Score drops more than AUROC	ratio comparison
H6	Statistical $\approx$ gradient boosting	$ \Delta \leq 0.02$

IV. Materials

Datasets:

- **Development:** MIMIC-IV v3.1 [8]
(65,078 adult ICU stays)
- **External validation:** eICU-CRD
v2.0 [9]
(200+ US ICUs, 7.9M person-hours)

Three Sepsis-3 label variants [1]:

Var.	Infection Window
A	Antibiotics and blood culture within 24 h of each other. SOFA baseline at ICU admission.
B	Antibiotics up to 72 h before culture, or 24 h after . SOFA baseline at ICU admission.
C	Antibiotics and blood culture within 72 h of each other. Lowest SOFA in previous 24 h.

In all three definitions, sepsis = suspected infection + SOFA score increase ≥ 2 points.

V. Methods: Model Architecture

The main model is a **dynamic discrete-time competing-risks hazard model** using **multinomial pooled logistic regression** [10].

- Predicts the **risk of sepsis every hour** during a patient's ICU stay.
- Does not require the proportional hazards assumption.
- Accounts for patients who **leave the ICU or die before developing sepsis** (competing risks).

Comparator Models

Gradient-Boosted Trees (GBT), Static Cox model, NEWS2, qSOFA, and SIRS.

V. Methods: Validation Strategy

Validation	What was done	Why?
Internal	Trained on patients from 2008–2016 and tested on 2017–2019 data [6].	To evaluate the model on newer, unseen patients.
Random split	Trained on 75% of patients and tested on 25%, chosen randomly [6].	To quantify how much the data split inflates reported performance (H4).
Pre-treatment	Used only data collected before the first antibiotic or blood culture [7].	To check whether the model is predicting sepsis or detecting when doctors start treatment.
External	Tested the same trained model on the eICU-CRD [9] dataset without retraining.	To see whether the model works in other hospitals.

Main evaluation metrics: Utility Score [11] and **Calibration**. AUROC was reported only for comparison because it may not detect clinically important performance failures [4].

V. Methods: Methodological Commitments

- **Pre-registered protocol:** The study plan was fixed before model development to reduce bias (following TRIPOD+AI guidelines [12]).
- **Leakage prevention:** The model only uses information available at each prediction time; missing values are handled using past information only.
- **Patient-level separation:** The same patient never appears in both training and testing data.
- **Frozen external validation:** The model trained on MIMIC-IV is directly tested on eICU-CRD without retraining.
- **Reproducible pipeline:** The complete workflow uses organised scripts, shared configurations, and publicly available code for replication.


VI. Results: Cohort and Label Variance

Variant	Name	ICU Stays	Sepsis Cases	%
A	Narrow	65,078	6,252	9.6%
B	Seymour-Standard	65,078	6,447	9.9%
C	Liberal	65,078	12,581	19.3%

The Liberal definition found almost **2× more sepsis cases** than the other two definitions.

The patient data and database were the same. Only the way sepsis was defined was different.

VI. Results: H1 Label Impact > Model Impact (CONFIRMED)



- **Label variation:** AUROC changed by **0.153** when the sepsis definition changed while using the same model.
- **Model variation:** AUROC changed by **0.145** when the model changed while using the same label.
- This confirms Cohen et al.'s [5] finding that sepsis label definition can strongly influence model performance.

The way sepsis is defined can affect performance as much as, or more than, the AI model itself.

VI. Results: H2 Coefficient Instability (CONFIRMED)

- **10 out of 25 features** changed their effect or strength significantly across different sepsis label definitions.
- Affected features include: MAP, temperature, bilirubin, platelets, vasopressors, norepinephrine/epinephrine, respiratory rate change, platelet measurement, WBC measurement, and P/F ratio.

Implication

A published hazard ratio for MAP could show a positive, negative, or almost no association with sepsis risk depending only on how the antibiotic-culture time window is defined [13].

Feature effects cannot be reliably interpreted without knowing which sepsis label definition was used.

VI. Results: H3 Treatment Leakage (CONFIRMED)

The treatment-anchored window [7] uses only information available **before the first antibiotic or blood culture**.

Metric	Value
Patient-hours analysed	179,003
Sepsis onset events	0

Critical Finding

No sepsis onset event is found before the clinical actions used to define the Sepsis-3 label.

This is not just a small drop in model performance (0.03 AUROC). The prediction target is **completely missing** in the clinically realistic time window.

Most MIMIC-IV sepsis prediction models are learning when doctors suspect sepsis and start treatment, rather than detecting when a patient's health is starting to get worse.

VI. Results: H4 Split Optimism (CONFIRMED)

Split	Model	AUROC	Utility Score
Temporal	Primary	0.607	-1.004
Random	Primary	0.730	-1.000
Temporal	GBT	0.746	-4.849
Random	GBT	0.720	-5.167

- **Primary model:** Random-split AUROC was **0.123 higher** than temporal-split AUROC (threshold: ≥ 0.02).
- The random split lets future patient patterns leak into training, inflating reported performance.
- This confirms Guo et al.'s [6] finding: **studies using random splits may overestimate real-world performance.**


VI. Results: H5 External Validation on eICU-CRD [9]

Variant	Model	AUROC (eICU)	AUROC (MIMIC)	Drop
A	Primary	0.736	0.760	0.024
B	Primary	0.437	0.607	0.170
C	Primary	0.725	0.759	0.034
n/a	NEWS2	0.605	0.63	0.025

- Variants A and C showed similar performance on both datasets, with only small drops.
- **Variant B (Seymour-Standard) showed a large performance drop of 0.170 AUROC** when tested on eICU-CRD.
- The most commonly used sepsis label definition did not transfer well to a different hospital dataset.

VI. Results: H6 Model-Class Difference (REJECTED)

Model-class effect on AUROC (Variant B, temporal split)
H6: $| \text{primary} - \text{GBT} | \leq 0.02$



- **GBT AUROC:** 0.746
- **Primary Model AUROC:** 0.607
- **Difference:** 0.139
- **GBT Utility Score:** -4.849
- **Primary Utility Score:** -1.004

The model ranking changes depending on the evaluation metric. A model that looks better by AUROC may perform worse in real clinical use.


VI. Results: Internal Performance (Temporal Test)

Variant	Model	AUROC	Utility Score	Alert Burden
A	Primary	0.760	-1.000	∞
A	GBT	0.742	-4.667	203
B	Primary	0.607	-1.004	∞
B	GBT	0.746	-4.849	194
C	Primary	0.759	-1.000	∞
C	GBT	0.752	-2.475	93
n/a	NEWS2	0.63–0.64	-1.1 to -1.6	97–223

All models had a negative Utility Score, meaning they did not provide clinical benefit and performed worse than giving no alerts.

Alert burden: For every real sepsis alert, the system generated **93–359 false alerts** (benchmark: 1.4 [14]).

VI. Results: Equity Analysis (Novel Finding)



- AUROC varied by up to **0.15** across racial subgroups (primary model, Variant B), from 0.56 (Other) to 0.71 (Asian-Chinese).
- The percentage of patients labelled as septic changed by **7.4–14.2 percentage points** across different label definitions within demographic subgroups.
- The **Unknown/Unable-to-obtain** race groups showed the largest changes (12.3–12.8 pp).
- This is the first study to examine both model performance disparities and label sensitivity across subgroups [15].

VI. Results: Variance Decomposition Summary

Factor	AUROC Change	Related Hypothesis	Status
Sepsis label definition	0.153	H1	CONFIRMED
Model choice	0.145	H6	REJECTED
Split method (temporal vs random)	0.123	H4	CONFIRMED
Evaluation metric (AUROC vs Utility Score)	0.170	H5	n/a

All four analyst decisions create AUROC changes of 0.12–0.17.

Most studies fix one sepsis label and compare models. This measures only one source of variation while ignoring three other sources of similar importance.

VII. Discussion: Key Finding

The performance collapse reported by Wang et al. [4] is not only a model evaluation issue. This study shows that it is strongly related to how the sepsis label is constructed.

The Sepsis-3 label in MIMIC-IV has three important limitations:

1. **Depends on clinical actions:** No sepsis cases are identified before treatment starts [7].
2. **Sensitive to label definition:** Different reasonable interpretations can produce almost double the number of sepsis cases [5].
3. **Produces poor clinical utility:** All tested models show negative Utility Scores [4].

These findings suggest that improving model complexity alone cannot solve the problem.
Better sepsis label design is required.

VII. Discussion: Contributions

#	Contribution
C1	Compared three Sepsis-3 label definitions using the same prediction model on MIMIC-IV [5].
C2	Showed no sepsis cases exist before antibiotics or culture tests, highlighting the label depends on clinical actions [7].
C3	Confirmed all models had negative Utility Scores, extending Wang et al. [4].
C4	Found 10 clinical features changed their effects across label definitions, supporting Lauritsen et al. [13].
C5	Quantified split optimism: random splits inflated AUROC by 0.123 vs. temporal splits, confirming Guo et al. [6].
C6	First study to examine how sepsis label definitions affect racial subgroup analysis [15].
C7	Used a pre-registered, TRIPOD+AI-compliant, fully reproducible pipeline [12].

VII. Discussion: Limitations

- **Small external event count:** The eICU dataset had only **63 sepsis cases**, making some evaluation metrics less reliable.
- **Deep learning comparison:** A deep survival model was planned but not implemented.
- **Limited features:** Clinical notes, procedure codes, and ventilator settings were not included.
- **Specific to Sepsis-3:** These findings apply to the Sepsis-3 label and may not apply to other sepsis definitions.

These limits are documented rather than hidden.

VIII. Conclusion and Future Work

1. **Better sepsis labels:** Define sepsis using objective patient measurements (such as lactate levels and organ function) instead of treatment decisions.
2. **Prospective study:** Test the new label in real clinical settings using future patient data.
3. **Better alert threshold:** Choose the alert threshold that maximises clinical benefit (Utility Score) instead of only improving prediction accuracy.
4. **Multi-hospital evaluation:** Test the model in different hospitals and identify an acceptable number of false alerts for clinical use.

References I

- [1] M. Singer, C. S. Deutschman, C. W. Seymour, et al., “The Third International Consensus Definitions for Sepsis and Septic Shock (Sepsis-3),” *JAMA*, vol. 315, no. 8, pp. 801–810, 2016. DOI: [10.1001/jama.2016.0287](https://doi.org/10.1001/jama.2016.0287).
- [2] K. E. Rudd, S. C. Johnson, K. M. Agesa, et al., “Global, regional, and national sepsis incidence and mortality, 1990–2017: Analysis for the Global Burden of Disease Study 2017,” *The Lancet*, vol. 395, no. 10219, pp. 200–211, 2020. DOI: [10.1016/S0140-6736\(19\)32989-7](https://doi.org/10.1016/S0140-6736(19)32989-7).
- [3] C. W. Seymour, F. Gesten, H. C. Prescott, et al., “Time to treatment and mortality during mandated emergency care for sepsis,” *New England Journal of Medicine*, vol. 376, no. 23, pp. 2235–2244, 2017. DOI: [10.1056/NEJMoa1703058](https://doi.org/10.1056/NEJMoa1703058).
- [4] Z. Wang, W. Wang, C. Sun, et al., “A methodological systematic review of validation and performance of sepsis real-time prediction models,” *npj Digital Medicine*, vol. 8, no. 1, p. 190, 2025.

References II

- [5] S. N. Cohen, J. Foster, P. Foster, et al., “Subtle variation in Sepsis-III definitions markedly influences predictive performance within and across methods,” *Scientific Reports*, vol. 14, no. 1, p. 1920, 2024. DOI: [10.1038/s41598-024-51762-5](https://doi.org/10.1038/s41598-024-51762-5).
- [6] L. L. Guo, S. R. Pfohl, J. Fries, et al., “Evaluation of domain generalization and adaptation on improving model robustness to temporal dataset shift in clinical medicine,” *Scientific Reports*, vol. 12, no. 1, p. 2726, 2022.
- [7] F. Kamran, D. Tjandra, A. Heiler, et al., “Evaluation of sepsis prediction models before onset of treatment,” *NEJM AI*, vol. 1, no. 3, 2024. DOI: [10.1056/AIoA2300003](https://doi.org/10.1056/AIoA2300003).
- [8] A. Johnson, L. Bulgarelli, L. Shen, et al., *MIMIC-IV, version 3.1*, PhysioNet, 2023. DOI: [10.13026/6mm1-ek67](https://doi.org/10.13026/6mm1-ek67).
- [9] T. J. Pollard, A. E. W. Johnson, J. D. Raffa, et al., *The eICU Collaborative Research Database, a freely available multi-centre database for critical care research*, 2018. DOI: [10.1038/sdata.2018.178](https://doi.org/10.1038/sdata.2018.178).

References III

- [10] R. B. D'Agostino, M.-L. Lee, A. J. Belanger, et al., "Relation of pooled logistic regression to time dependent Cox regression analysis: The Framingham Heart Study," *Statistics in Medicine*, vol. 9, no. 12, pp. 1501–1515, 1990. DOI: 10.1002/sim.4780091214.
- [11] M. A. Reyna, C. S. Josef, R. Jeter, et al., "Early prediction of sepsis from clinical data: The PhysioNet/Computing in Cardiology Challenge 2019," *Critical Care Medicine*, vol. 48, no. 2, pp. 210–217, 2020. DOI: 10.1097/CCM.0000000000004145.
- [12] G. S. Collins, K. G. M. Moons, P. Dhiman, et al., "TRIPOD+AI statement: Updated guidance for reporting clinical prediction models that use regression or machine learning methods," *BMJ*, vol. 385, e078378, 2024. DOI: 10.1136/bmj-2023-078378.
- [13] S. M. Lauritsen, B. Thiesson, M. J. Jørgensen, et al., "The framing of machine learning risk prediction models illustrated by evaluation of sepsis in general wards," *npj Digital Medicine*, vol. 4, no. 1, p. 158, 2021. DOI: 10.1038/s41746-021-00529-x.


References IV

- [14] M. Moor, N. Bennett, D. Plečko, et al., “Predicting sepsis using deep learning across international sites: A retrospective development and validation study,” *eClinicalMedicine*, vol. 62, p. 102 124, 2023.
- [15] H. Wang, Y. Li, A. Naidech, and Y. Luo, “Comparison between machine learning methods for mortality prediction for sepsis patients with different social determinants,” *BMC Medical Informatics and Decision Making*, vol. 22, no. Suppl 2, p. 156, 2022.

Appendix: Label Sensitivity vs. AUROC

Label sensitivity vs. AUROC by subgroup

Novel conjunction: Wang et al. (2022) established each separately



Thank You

Questions & Discussion

Abdul Ahad

ahadarif.1998@gmail.com

Atlantic Technological University (ATU), Galway